

# Report

## Word-Level Handwritten OCR using IAM Dataset

### Team Members

Baris Surmelioglu – 152048

Ozan Polat – 150918

### 1. Problem Description

The goal of this project is to solve a **word-level handwritten Optical Character Recognition (OCR)** problem. Given an image containing a **single handwritten word**, the task is to predict its **textual transcription**. This is a challenging problem due to large variations in handwriting styles, character spacing, word length, and image quality.

### 2. Dataset Description

We use the IAM Handwriting Database (word-level), accessed via a publicly available **Kaggle dataset**. The dataset provides handwritten word images along with ground-truth transcriptions and follows the original IAM annotation format.

the handwritten word structure.

- Annotation file: words\_new.txt
- Only samples with status = ok are used
- Missing or corrupted images are automatically filtered out
- The dataset is limited to word-image folders from **A to E**, as later folders are not visually available

### Dataset Split

The dataset is split using CSV files:

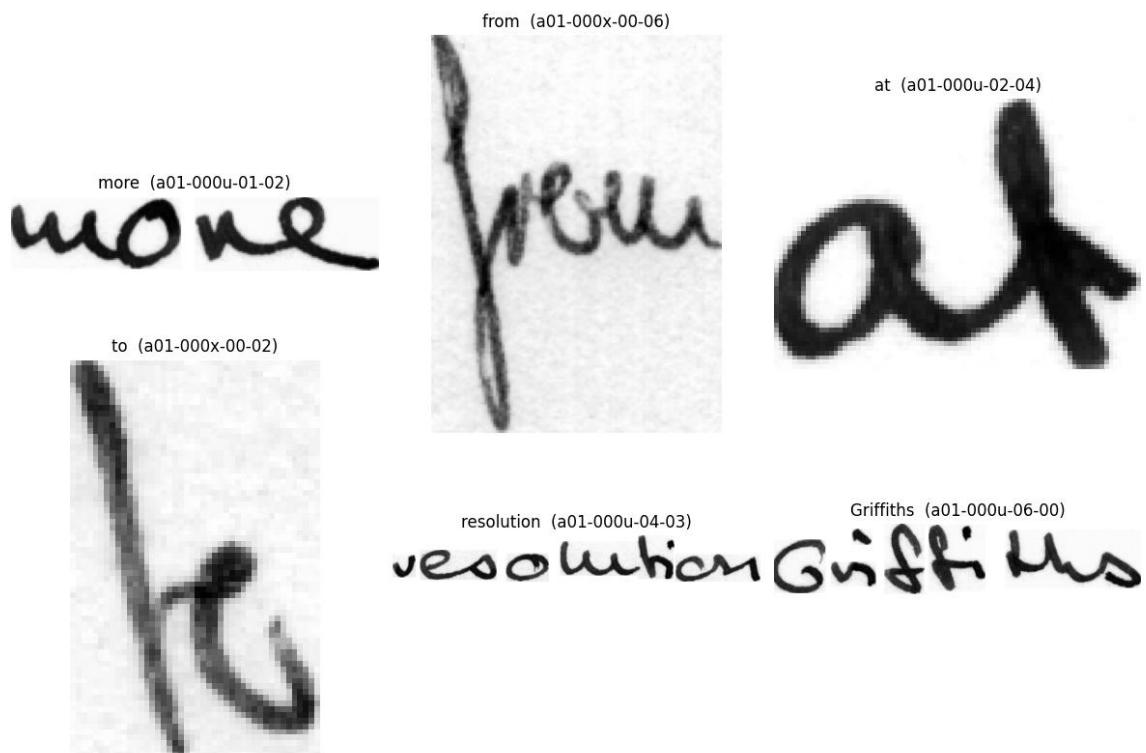
- splits/train.csv
- splits/val.csv
- splits/test.csv

Split ratios:

- **Train:** 80%
- **Validation:** 10%
- **Test:** 10%

The final **test set contains 3825 word images.**

Samples:



### 3. Preprocessing

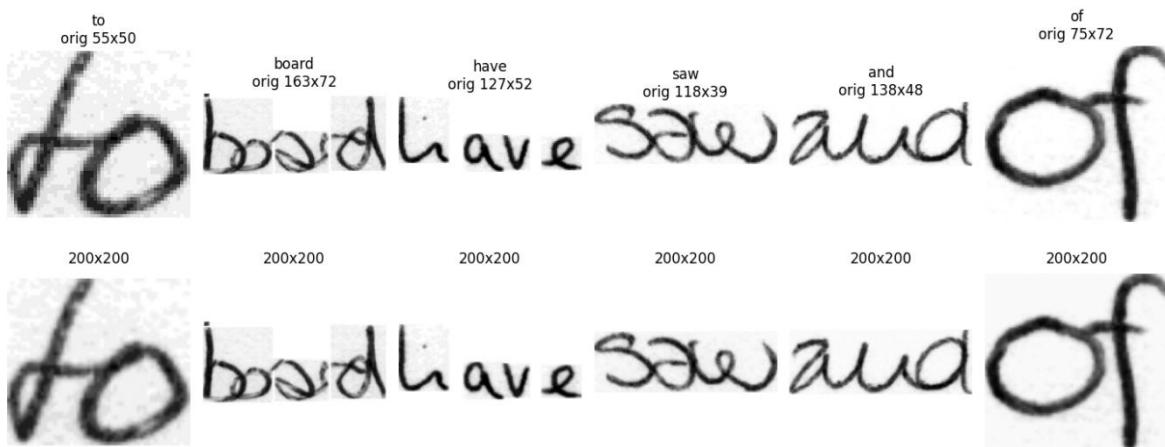
Each word image is processed as follows:

- Converted to **grayscale**
- Resized to **200×200 pixels**
- Aspect ratio preserved, padded and centered
- Pixel values normalized to [0, 1]

## Data Augmentation

To improve generalization, light augmentation is applied **only to the training set**:

- For data augmentation, RandomAffine transformations with small rotation and translation ranges were applied (e.g., rotations up to  $\pm 5$  degrees and small translations), in order to increase robustness without distorting



## 4. Model Architecture

The OCR system follows a **CNN–RNN–CTC** pipeline.

### CNN Backbone

- **ResNet18**, pretrained on ImageNet
- First convolution layer adapted from 3-channel to **1-channel input**

### Sequence Modeling

- Feature maps are converted into sequences
- **2-layer Bidirectional LSTM** is used to model character dependencies

### Output

- Linear projection layer to character logits
- **CTC Loss** for alignment-free training
- **Greedy decoding** during inference

### Architecture flow:

ResNet18 → Feature Maps → Sequence Extraction → BiLSTM → Linear → CTC

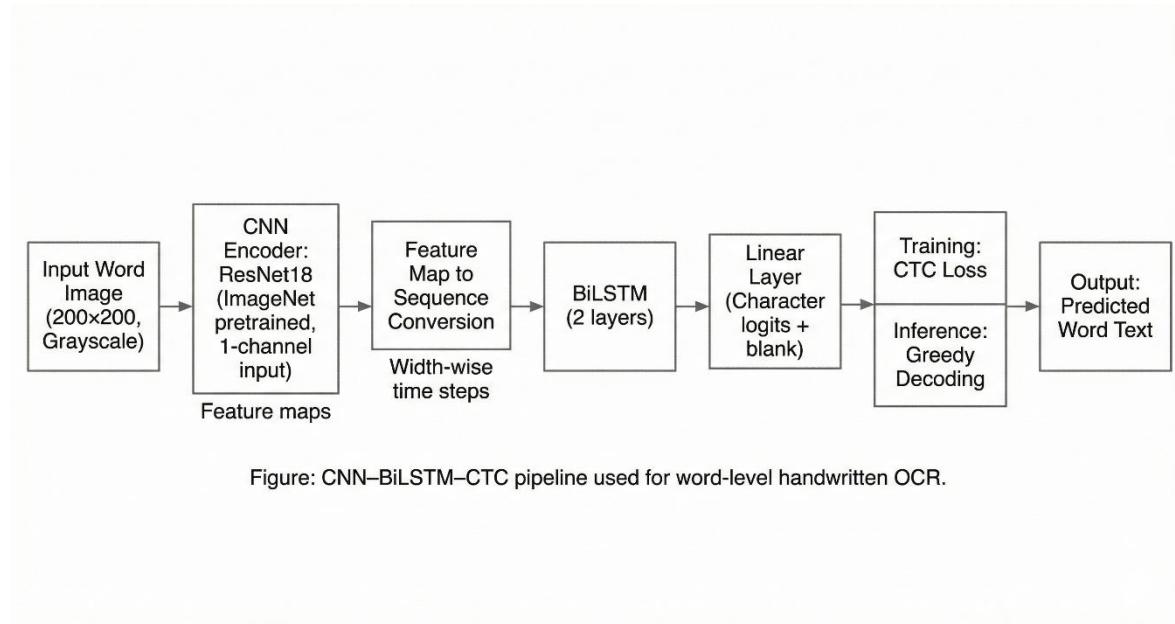


Figure: CNN–BiLSTM–CTC pipeline used for word-level handwritten OCR.

## 5. Model Analysis

- Backbone: ResNet18
- Sequence model: 2× BiLSTM
- Total parameters: **16 million**:
  - Model file size: **61 MB**
  - The model size fits comfortably in GPU memory
  - Number of parameters is dominated by the CNN backbone
- Suitable for word-level OCR without excessive computational cost

## 6. Training Procedure

Training is performed using the following setup:

- **Optimizer:** AdamW, chosen for its better regularization properties compared to Adam
- **Learning rate:** 1e-4
- **Scheduler:** ReduceLROnPlateau
- **Batch size:** 32


- **Gradient clipping** enabled to stabilize training
- Best model selected based on **lowest validation CER**

Training and validation metrics are logged per epoch.

## Commands to Run

The training and evaluation pipeline can be executed by running the notebook:

*IAM\_Word\_Level\_OCR.ipynb*



Training complete. Best CER: 0.1079

## 7. Hyperparameters and Explanation

Hyperparameter	Value	Explanation
Image size	200×200	Standardized input size while preserving details
Batch size	32	Balanced between stability and memory usage
Learning rate	1e-4	Stable convergence for pretrained backbone
Optimizer	AdamW	Better regularization than Adam
Scheduler	ReduceLROnPlateau	Adapts learning rate based on validation CER
Epochs	10	Enough to converge without overfitting

## 8. Loss and Evaluation Metrics

### Loss

- **CTC Loss** is used to handle variable-length predictions without explicit alignment.

## Metrics

- **CER (Character Error Rate)**
- **WER (Word Error Rate)**
- **Word Accuracy**

At least two metrics (CER and WER) are used as required.

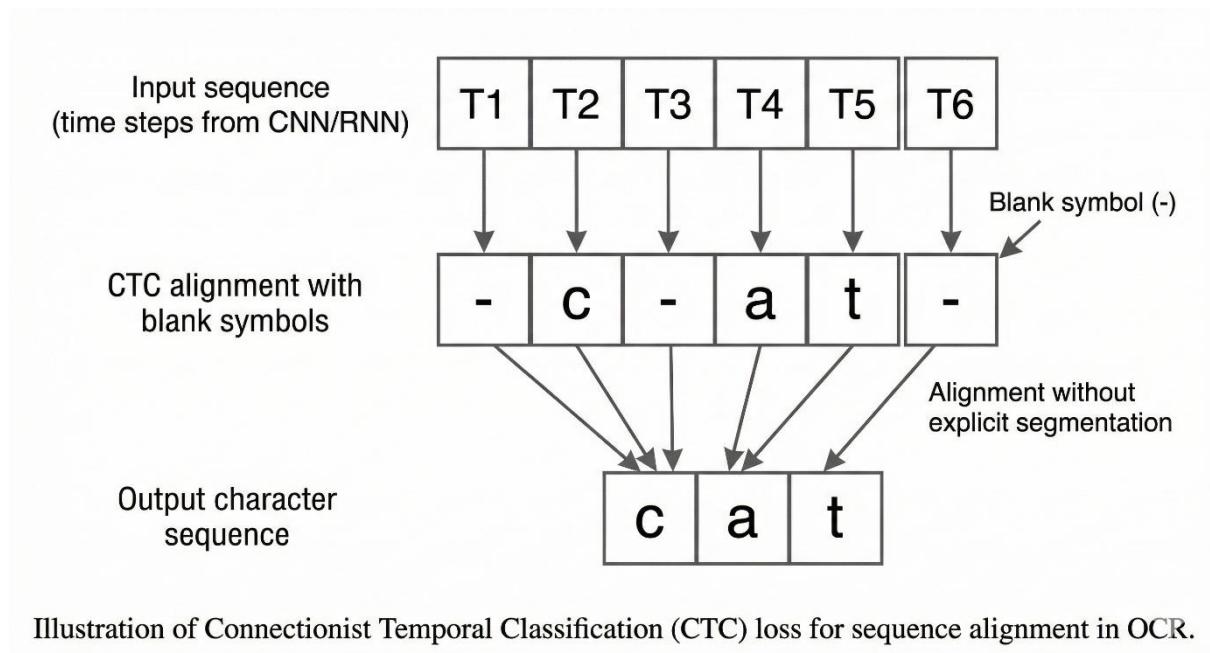


Illustration of Connectionist Temporal Classification (CTC) loss for sequence alignment in OCR.

## 9. Results

### Test Set Results (Best Model)

- **Test CER:** 0.1166
- **Test WER:** 0.3354
- **Test Word Accuracy:** 0.6646
- **Test samples:** 3825

The model shows consistent performance on both validation and test sets.

Errors mostly occur on longer words or visually ambiguous characters.

GT: the  
Pred: the  
(b04-054-03-04)

the Berlin NATO

GT: Berlin  
Pred: Berlin  
(a05-089-07-02)

GT: NATO  
Pred: ITO  
(a02-020-05-08)

GT: elegant  
Pred: elegant  
(e01-059-04-04)

GT: yourself  
Pred: yarselv  
(d04-053-09-01)

GT: spite  
Pred: spite  
(b05-088-01-07)

elegant yourself spite

GT: the  
Pred: the  
(c03-087d-04-01)

GT: more  
Pred: more  
(c03-087c-06-01)

GT: city  
Pred: city  
(c02-030-02-04)

the more city

GT: of  
Pred: of  
(b05-062-03-09)

GT: Banking  
Pred: Baking  
(a01-049x-04-04)

GT: to  
Pred: to  
(a03-011-04-05)

of Banking to

## **10. Training and Inference Time**

- Training time: ~30 minutes for 10 epochs on NVIDIA GPU
- Inference time: ~5ms per image (batch size 32)

## **11. Comparison of Models**

Only one main architecture is used in this project.

The choice of ResNet18 + BiLSTM provides a good balance between accuracy and computational cost for word-level OCR.

## **12. Libraries and Tools**

A full list of libraries is provided in requirements.txt, including:

- PyTorch
- torchvision
- NumPy
- matplotlib
- Pillow (PIL)
- TensorBoard

## **13. TensorBoard Experiment Tracking**



Training metrics were logged using TensorBoard for monitoring loss and error rates.



### **Logged Metrics:**

- Training loss (per batch and per epoch)
- Validation loss
- Validation CER and WER
- Learning rate changes

### **Analysis:**

- Training loss decreased steadily over epochs
- Validation CER dropped from ~0.45 to ~0.11, showing good convergence
- Learning rate was reduced automatically when validation loss plateaued
- No significant overfitting observed (validation loss follows training loss)





## 14. Runtime Environment

- Python 3
- GPU-enabled environment recommended
- Experiments conducted using standard deep learning libraries

## **15. Completed Items Table**

OCR of handwritten text	2 point
pre-trained model on the different problem (transfer-learning)	1 point
Adaptive hyperparameters	1 point
Data augmentation	1 point
Tensorboard	1 point

## **16. GitHub Repository**

Link to GitHub repository:

<https://github.com/BarisAI/Computer-Vision-OCR-of-handwritten-text>

## **17. Bibliography**

- Marti, U.-V., & Bunke, H. “The IAM-database: an English sentence database for offline handwriting recognition.”
- Graves, A. et al. “Connectionist Temporal Classification.”
- He, K. et al. “Deep Residual Learning for Image Recognition.”
- Kaggle Dataset:

IAM Handwriting Word Database. Available at:

<https://www.kaggle.com/datasets/nibinv23/iam-handwriting-word-database>